

# ■ Machine Learning Roadmap (6 Months)

This roadmap provides a structured 6-month learning plan to become proficient in Machine Learning. Each month covers focused topics with detailed descriptions, explanations, and outcomes.

## ➡■ Month 1 — Foundations (Math + Python)

**Python for Data Science:** Learn the essential Python libraries for ML: NumPy (numerical arrays & math), Pandas (data manipulation), and Matplotlib/Seaborn (data visualization). These tools will help you prepare, analyze, and visualize datasets. Outcome: You can clean, preprocess, and explore any dataset efficiently.

**Mathematics Refresher:** Build strong foundations in the math that powers ML models. Focus on: Linear Algebra (vectors, matrices, dot products), Probability (distributions, Bayes theorem), and Calculus (derivatives, partial derivatives, optimization). Outcome: You'll understand the core math behind ML algorithms.

**Mini Projects:** Apply Python and Math knowledge: explore Titanic dataset survival rates with Pandas and Seaborn, simulate probability experiments (like dice rolls). Outcome: Practical confidence with tools and math basics.

## ➡■ Month 2 — Core ML Concepts

**Supervised Learning:** Dive into regression and classification models. Learn Linear Regression for predicting numerical values and Logistic Regression for binary classification problems. Explore concepts like cost functions and gradient descent. Outcome: You can train, evaluate, and interpret simple predictive models.

**Unsupervised Learning:** Understand clustering (KMeans, DBSCAN) and dimensionality reduction (PCA). These techniques help find hidden patterns and simplify high-dimensional data. Outcome: You'll know how to group unlabeled data and extract meaningful structure.

**Projects:** Predict house prices using regression. Cluster customer segments with KMeans. Outcome: Experience applying ML to real-world datasets.

## ➡■ Month 3 — Practical ML & Scikit-Learn

**ML Workflows:** Learn how to structure end-to-end ML projects. Cover train/test splits, cross-validation, pipelines, and preprocessing (feature scaling, categorical encoding). Outcome: You can build reusable and clean ML workflows.

**Model Evaluation:** Study evaluation metrics: Accuracy, Precision, Recall, F1, ROC curves. Learn about overfitting, bias-variance tradeoff, and regularization techniques (L1/L2). Outcome: You'll know how to measure model quality and avoid common pitfalls.

**Projects:** Build a Spam Email Classifier and Titanic survival prediction model using pipelines and metrics. Outcome: Practical applied ML experience.

## ➡■ Month 4 — Neural Networks & Deep Learning

**Neural Network Basics:** Understand perceptrons, activation functions (sigmoid, ReLU), forward propagation, and backpropagation. Manually build a neural network with NumPy to grasp fundamentals. Outcome: A solid mental model of neural nets.

**Deep Learning Frameworks:** Learn TensorFlow or PyTorch for real-world neural networks. Train a Multi-Layer Perceptron on MNIST digits. Explore optimizers (SGD, Adam), dropout, and batch normalization. Outcome: You can train your first neural networks.

**Project:** Handwritten digit recognizer (MNIST). Outcome: A complete deep learning project under your belt.

## ➡■ Month 5 — Specialized Topics

**Natural Language Processing (NLP):** Work with text: Bag of Words, TF-IDF, Word embeddings (Word2Vec, GloVe). Outcome: You can preprocess and classify text data.

**Computer Vision (CV):** Learn Convolutional Neural Networks (CNNs) for image recognition tasks. Outcome: Ability to build image classifiers.

**Time Series Forecasting:** Explore ARIMA models and recurrent neural networks (RNNs). Outcome: Skills to forecast temporal data like sales or stock prices.

**Recommendation Systems:** Build recommenders using collaborative filtering and matrix factorization. Outcome: Understanding of personalization systems.

## ➡■ Month 6 — Advanced & Real-World ML

**MLOps Basics:** Learn model deployment using Flask/FastAPI, containerization with Docker, and model versioning practices. Outcome: You can serve ML models as APIs in production.

**Advanced ML:** Explore powerful algorithms like XGBoost, LightGBM, and transfer learning using pre-trained models (ResNet, BERT). Outcome: You can tackle state-of-the-art ML problems.

**Capstone Project:** Choose a project: build an end-to-end ML pipeline or compete in a Kaggle competition. Outcome: Portfolio-ready project demonstrating your ML expertise.

■ **End State:** By the end of 6 months, you'll be confident with ML foundations, practical tools, and have portfolio-ready projects to showcase your skills.